

# Chapter 29


## Multi-threading

# Introduction

- Since Java 1, one of the biggest draws to Java was the internal support for multi-threading applications.
- In other languages (like C) you need to import system dependent libraries in order to have threading inside of a program.
- Threading is normally taken care of at the OS kernel level, but in Java we have direct access to non system dependent methods for threading.

# Thread Concepts


- Threading is the act of running a piece of your program concurrently to the rest of your program.
- A thread can be thought of as the flow of execution for a single task.



The flow of three threads, running in a multi-core / multi-processor environment

# Thread Concepts

- You don't need to have multiple processors or multiple cores to get a benefit from running multiple threads.



One possible execution for a single processor, single core computer running a multi-threading program. Notice that the order of executions is NOT guaranteed This is important!

# Creating a “Runnable” thread

- The first thing that you must realize is that in Java, threading is done on a class level.
- Any class can be allowed to run in its own thread by declaring that the thread implements Runnable.
- The `java.lang.Runnable` interface has only a single method that must be implemented.
- That method is `public void run()`
- The `run()` method returns nothing, and throws no checked exceptions!

# Creating a thread

```
public class PrintChars implements Runnable{
    private char c;
    private int times;

    public PrintChars(){
        this('a', 100);
    }

    public PrintChars(char charToPrint, int times){
        this.c = charToPrint;
        this.times = times;
    }

    public void run(){
        for (int i = 0; i < times; i++){
            System.out.print(c);
        }
    }
}
```

# Creating a thread

```
public class PrintNumbers implements Runnable{
    private int start;
    private int range;

    public PrintNumbers(){
        this(0, 100);
    }

    public PrintNumbers(int start, int range){
        this.start = start;
        this.range = range;
    }

    public void run(){
        for (int i = start; i < range; i++){
            System.out.print(i);
        }
    }
}
```

# Creating a thread

```
public class PrintCharsMain{
    public static void main(String[] args){

        PrintChars printA = new PrintChars();
        PrintChars printB = new PrintChars('c', 100);
        PrintChars printC = new PrintChars('d', 100);

        PrintNumbers printNum = new PrintNumbers();

        Thread t1 = new Thread(printA);
        Thread t2 = new Thread(printB);
        Thread t3 = new Thread(printC);
        Thread t4 = new Thread(printNum);

        t1.start();
        t2.start();
        t3.start();
        t4.start();

        System.out.println("End of main thread");
    }
}
```





# Breakdown

- You can use the following steps to create, and execute a thread.
  - 1) Create a class that implements `java.lang.Runnable`.
  - 2) Implement the `run` method as if it were the main method of that class.
  - 3) Create a `Thread` object and use the class you declared `Runnable` as an argument to the constructor.
  - 4) Call the `Thread.start()` method on that thread object.
  - 5) NOTE: calling the `run()` method directly will NOT produce a new thread, it will simply run that method in the current thread.

# The Thread class

- The thread class contains all of the methods for controlling threads.
- The thread class also implements Runnable, so you can extend Thread and make a class that can run itself in a new thread. This is generally considered bad practice and should be avoided.

## java.lang.Thread

```
Thread()  
Thread(task: Runnable)  
start(): void  
isAlive(): boolean  
setPriority(p: int): void  
join(): void  
sleep(millis: long): void  
yield(): void  
interrupt(): void
```

# The Thread class

- The Thread class has two static methods that are very useful for ensuring that all threads have a chance to run.
- The `yield()` and `sleep()` methods allow the currently running (calling) method to step aside and allow other waiting threads to run.
- This becomes necessary only when dealing with threads of different priority, as high priority threads may starve a low priority thread.

# yield() and sleep()

```
public class PrintChars implements Runnable{
    private char c;
    private int times;

    public PrintChars(){
        this('a', 100);
    }

    public PrintChars(char charToPrint, int times){
        this.c = charToPrint;
        this.times = times;
    }

    public void run(){
        for (int i = 0; i < times; i++){
            System.out.print(c);
            if (i % 10 == 0 && i != 0){
                Thread.yield();
            }
        }
    }
}
```

This new addition to the PrintChars class will allow a waiting thread to execute every 10 iterations.

# yield() and sleep()

```
File Edit View Terminal Help
james@renegade$ java PrintCharsMain
aaaaaaaaaaaaEnd of main thread
dcdadaaaaaaaaaa0aaaaaaaaaadddddddcdccccccccdaaaaaaaaaa1aaaaaaaaaadddddddcccc
ccccdddaaaaaaaaaa2aaaaaaaaaadddddddcccccccccccccccccccccccccccccccccccccccc
aaaaaaaaaaaaaaaaaddddccccccccccccdd4ddcccccccccccc5dddddccccdd6ddcccccccccccc
cccccccccccccccccd7ddcccccd8cccccccccccc9cccccccccccc10cccccccc11cc12131415161718192
02122232425262728293031323334353637383940414243444546474849505152535455565758596
0616263646566676869707172737475767778798081828384858687888990919293949596979899j
ames@renegade$
```

We can see some of the effects of the yield call, They will be even more prominent in a single CPU system. Below is what happens if we yield after printing every character.

```
File Edit View Terminal Help
james@renegade$ java PrintCharsMain
aEnd of main thread
accdca0ccddccc1a2ccddcc3a4ccddcc5a6ccddcc7a8ccddcc9a10ccddcc11a12ccddcc13a14ccddcc15a1
6ccddccc17a18ccddcc19a20ccddccc21a22ccddccc23a24ccddccc25a26ccddcc27a28ccddcc29a30ccddcc
31a32ccddcc33a34ccddcc35a36ccddccc37a38ccddccc39a40ccddcc41a42ccddcc43aa44d45a46d47a48
d49aaaa50d51aaaaaaaa52d53aaaa54d55aaaaaaaa56d57aaaa58d59aaa60d61aaa62d63a64d65aa
aa66d67aaaaaaaa68d69aaaaa70d71aaaa72d73aaaaa74d75aaa76d77aaaa78d79aaaaa80d81d82d8
3d84d85d86d87d88d89d90d91d92d93d94d95d96d97d98d99dddddcccccccccccccccccccccccccccc
dddddjames@renegade$
```

# yield() and sleep()

- The sleep method does exactly what it sounds like it should, it puts the thread to sleep for some number of milliseconds.
- After the time as elapsed, then the thread resumes and begins executing code again.
- You can use the Thread.sleep() for a number of reasons, everything from waiting for a resource to become free to dramatic effect in your output.

# yield() and sleep()

```
public class PrintChars implements Runnable{
    private char c;
    private int times;

    public PrintChars(){
        this('a', 100);
    }

    public PrintChars(char charToPrint, int times){
        this.c = charToPrint;
        this.times = times;
    }

    public void run(){
        for (int i = 0; i < times; i++){
            System.out.print(c);
            if (i % 10 == 0 && i != 0){
                try{
                    Thread.sleep(100);
                }
                catch (InterruptedException ignore){
                }
            }
        }
    }
}
```



# yield() and sleep()

[illegible]

As we can see from the output, after 10 characters, the thread sleeps and other threads have the ability to take over. This is different from `yield` in that if you call `yield`, you are not guaranteeing that another thread will get to run. `Sleep` will always allow another thread to run (if there is one to run)

# Thread Priority

- You can control which thread has the most runtime by changing its priority.
- The priorities are arranged from 1 – 10 with 10 being the highest priority.
- Although you can use the integer numbers directly, you should use the predefined constants
  - 1) `Thread.MAX_PRIORITY` = 10
  - 2) `Thread.NORM_PRIORITY` = 5
  - 3) `Thread.MIN_PRIORITY` = 1

# Thread Priority

- The JVM will always pick the highest priority thread that is ready to run.

```
Thread t1 = new Thread(printA);
Thread t2 = new Thread(printB);
Thread t3 = new Thread(printC);
Thread t4 = new Thread(printNum);

t1.setPriority(Thread.MAX_PRIORITY);

t1.start();
t2.start();
.
.
```

[illegible]

# Thread Priority

- When setting different priorities, this is where the `yield()` and `sleep()` methods are necessary to make sure that high priority threads don't starve the lower priority threads.
- This is easy to do in systems where you have only a single processor, or where you are running more threads than processors and time sharing is necessary.

# Joining Threads

- What if only a piece of your task is able to be made concurrent?
- What if you need to guarantee that all of the operations from some other threads are complete before the program can continue?
- We have the ability to call `Thread.join()`
- This call will cause the current thread to block until the thread that was `join()`'d finishes.

# Joining Threads

- We have noticed in the last few slides that our program prints that the main thread is done at random times of execution.
- We don't want this behavior, so we will join all of the printing threads to the main thread.
- Doing this will cause the main thread to wait before finishing its execution.

# Joining Threads

```
public class PrintCharsMain{
    public static void main(String[] args){

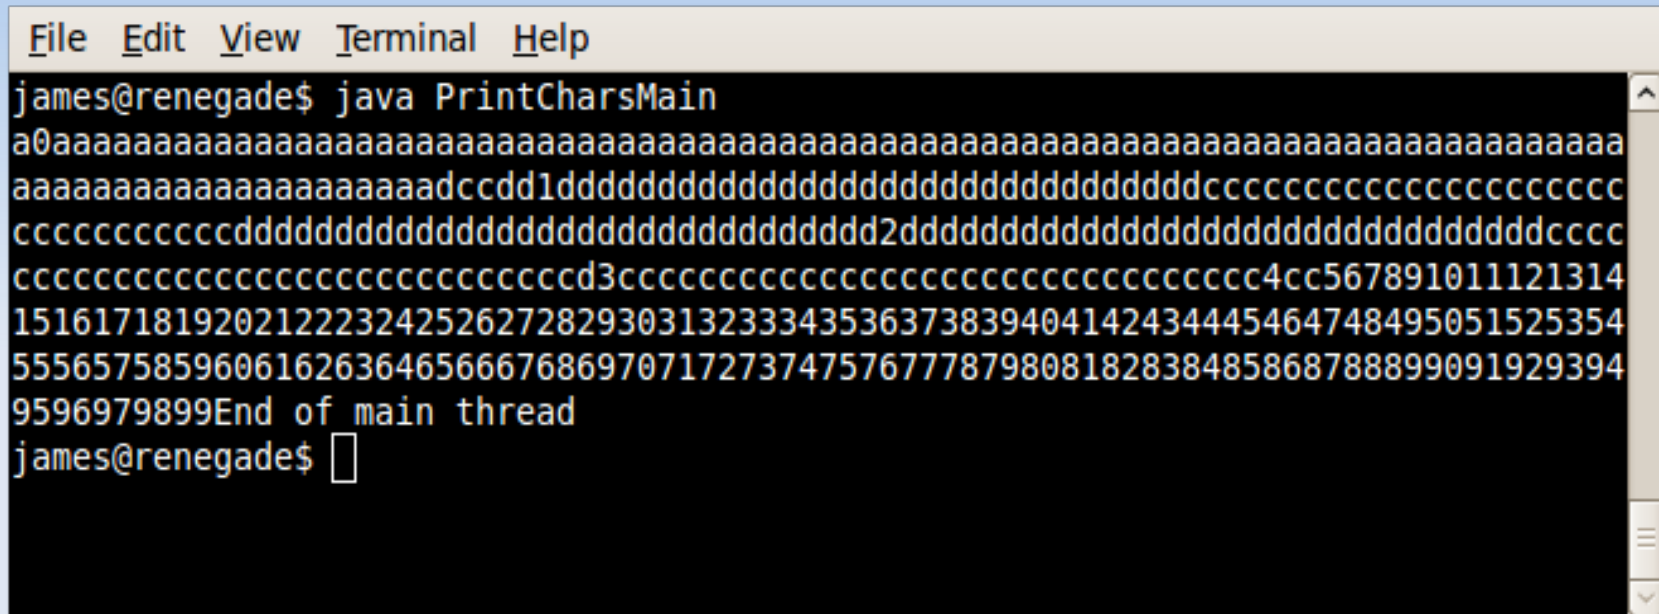
        PrintChars printA = new PrintChars();
        PrintChars printB = new PrintChars('c', 100);
        PrintChars printC = new PrintChars('d', 100);
        PrintNumbers printNum = new PrintNumbers();

        Thread t1 = new Thread(printA);
        Thread t2 = new Thread(printB);
        Thread t3 = new Thread(printC);
        Thread t4 = new Thread(printNum);

        t1.start();
        t2.start();
        t3.start();
        t4.start();
        try{
            t1.join();
            t2.join();
            t3.join();
            t4.join();
        }catch(InterruptedException ignore){}

        System.out.println("End of main thread");
    }
}
```

# Joining Threads

A terminal window with a menu bar (File, Edit, View, Terminal, Help) and a black background. The text is white. It shows the execution of a Java program. The output consists of a long string of 'a's, followed by a sequence of 'd' and 'c' characters, then a sequence of numbers from 1 to 95, and finally the text "End of main thread". The prompt "james@renegade\$" is visible at the bottom.

```
File Edit View Terminal Help
james@renegade$ java PrintCharsMain
a0aaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaa
aaaaaaaaaaaaaaaaaaaaadccdd1ddddddddddddddddddddddddddddddcccccccccccccccccccc
cccccccccccccccccccccccccccccccccccccccccccccccccccccccccccccccccccccccccccc
cccccccccccccccccccccccccccccccccccccccccccccccccccccccccccccccccccccccccccc
15161718192021222324252627282930313233343536373839404142434445464748495051525354
55565758596061626364656667686970717273747576777879808182838485868788899091929394
9596979899End of main thread
james@renegade$
```

No matter how many times we run this program, it will always end with “End of main thread”.



# Thread Synchronization

- What happens when a single resource is shared by multiple threads?
- If all of the threads are simply reading, then this is not a big problem, but if each thread is reading and writing then we can get corrupt or stale data.
- We need ways to force two or more threads to play nice and wait until one thread is completely finished with a resource before using it.

# Thread Synchronization

- There are several ways to synchronize threads, but first let's look at an example of why we may need to synchronize two threads that are performing tasks on a shared resource.
- This is a famous example that is in every thread text book, including the one in this class.
- It is the husband and wife at the ATM problem.
- Imagine a husband and wife at two separate ATM machines performing transactions on the same account unaware of the other person.

# Thread Synchronization

```
public class Account{  
    private int balance = 0;  
  
    public int getBalance(){  
        return balance;  
    }  
  
    public void deposit(int amount){  
        balance = balance + amount;  
    }  
}
```

# Thread Synchronization

```
public class UseAccount{
    private static Account account = new Account();

    public static void main(String[] args){
        Thread husband = new Thread(new Person());
        Thread wife = new Thread(new Person());

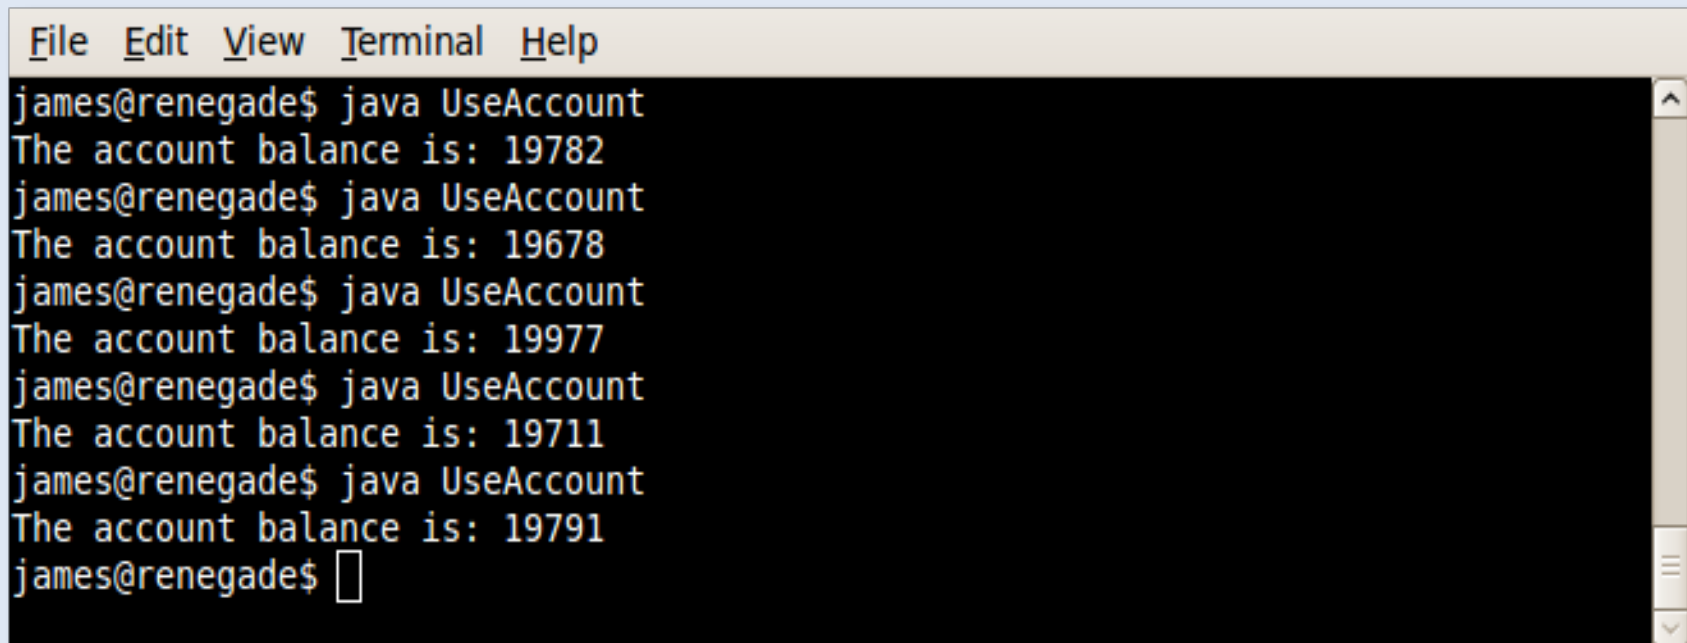
        husband.start();
        wife.start();
        try{
            husband.join();
            wife.join();
        }catch(InterruptedException ignore){}

        System.out.println("The account balance is: " +
                           account.getBalance());
    }

    static class Person implements Runnable{
        public void run(){
            for (int i = 0; i < 10000; i++){
                account.deposit(1);
            }
        }
    }
}
```

# Thread Synchronization

- Although the previous example looks okay at first, and it should say we have \$20,000 in the account at the end, we end up with random and corrupt results.

A terminal window with a menu bar (File, Edit, View, Terminal, Help) and a black background. It shows a series of commands and outputs from a user named 'james' on a machine named 'renegade'. The user runs 'java UseAccount' five times, each time receiving a different account balance: 19782, 19678, 19977, 19711, and 19791. The terminal window has a scrollbar on the right side.

```
File Edit View Terminal Help
james@renegade$ java UseAccount
The account balance is: 19782
james@renegade$ java UseAccount
The account balance is: 19678
james@renegade$ java UseAccount
The account balance is: 19977
james@renegade$ java UseAccount
The account balance is: 19711
james@renegade$ java UseAccount
The account balance is: 19791
james@renegade$
```

# Thread Synchronization

- So what when wrong? What we didn't realize is that the statement  $\text{balance} = \text{balance} + \text{amount}$  is broken up into two main parts.
  - 1)  $\text{Balance} + \text{amount}$
  - 2)  $\text{Balance} = \text{that number}$

<u>Step</u>	<u>balance</u>	<u>Thread 1</u>	<u>Thread 2</u>
1	0	$\text{balance} + \text{amount}$	
2	0		$\text{balance} + \text{amount}$
3	1	$\text{balance} = \text{new amount}$	
4	1		$\text{balance} = \text{new amount}$

Because the balance global variable doesn't change until after the assignment statement has been executed, the second thread reads stale data. This is what is known as a race condition.

# Thread Synchronization

- So how do we fix this?
- One thing we can do is use the `synchronized` keyword to synchronize the method.

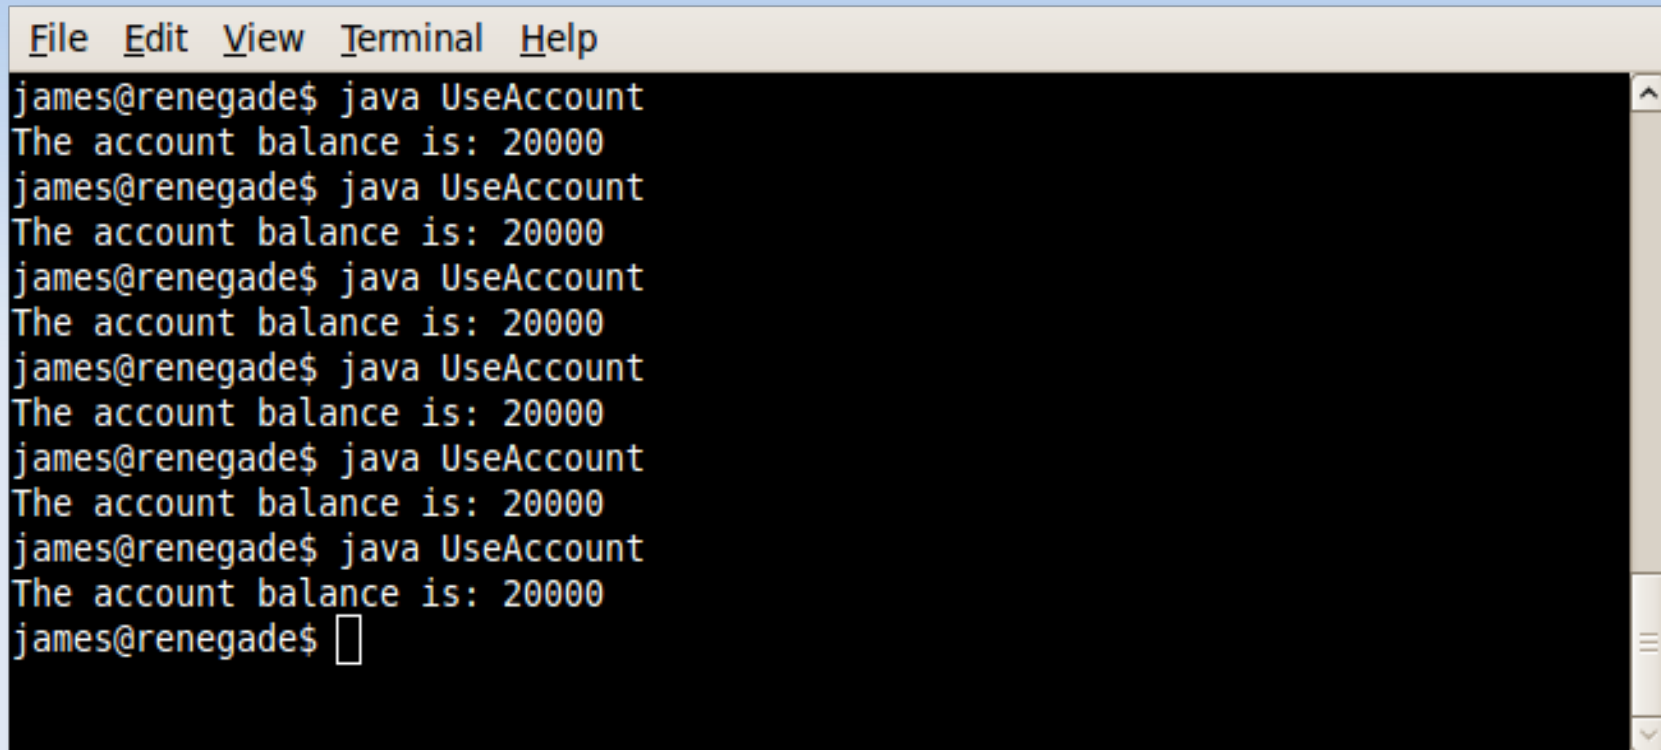
```
public class Account{  
    private int balance = 0;  
  
    public synchronized int getBalance(){  
        return balance;  
    }  
  
    public synchronized void deposit(int amount){  
        balance = balance + amount;  
    }  
}
```

# Thread Synchronization

- What this does is allow the first thread that calls that method to obtain a lock on the account object.
- Remember that when a thread enters a synchronized method that it locks the WHOLE object. This means that any other synchronized methods that are in the class will also be locked out until it has finished executing.



# Thread Synchronization

A terminal window with a menu bar (File, Edit, View, Terminal, Help) and a black background. It shows a series of commands and outputs. The command 'java UseAccount' is entered six times, each followed by the output 'The account balance is: 20000'. The prompt 'james@renegade\$' is visible at the start of each line. The window has a scrollbar on the right side.

```
File Edit View Terminal Help
james@renegade$ java UseAccount
The account balance is: 20000
james@renegade$ java UseAccount
The account balance is: 20000
james@renegade$ java UseAccount
The account balance is: 20000
james@renegade$ java UseAccount
The account balance is: 20000
james@renegade$ java UseAccount
The account balance is: 20000
james@renegade$ java UseAccount
The account balance is: 20000
james@renegade$
```

No matter how many times we run this class, we will see the same \$20,000 in the final balance. This is because one thread is forced to wait until another has completed its operations. This alleviates the race condition.

# Thread Synchronization

- What happens if you don't want to synchronize an entire method because that would cause wasteful blocking?
- We can use the synchronized keyword to lock any object in a synchronized block
- This can be useful to allow a long method to only lock what is absolutely necessary.

# Thread Synchronization

```
public class UseAccount{
    private static Account account = new Account();

    public static void main(String[] args){
        Thread husband = new Thread(new Person());
        Thread wife = new Thread(new Person());

        husband.start();
        wife.start();
        try{
            husband.join();
            wife.join();
        }catch(InterruptedException ignore){}

        System.out.println("The account balance is: " + account.getBalance());
    }

    static class Person implements Runnable{
        public void run(){
            for (int i = 0; i < 10000; i++){
                synchronized(account){
                    account.deposit(1);
                }
            }
        }
    }
}
```

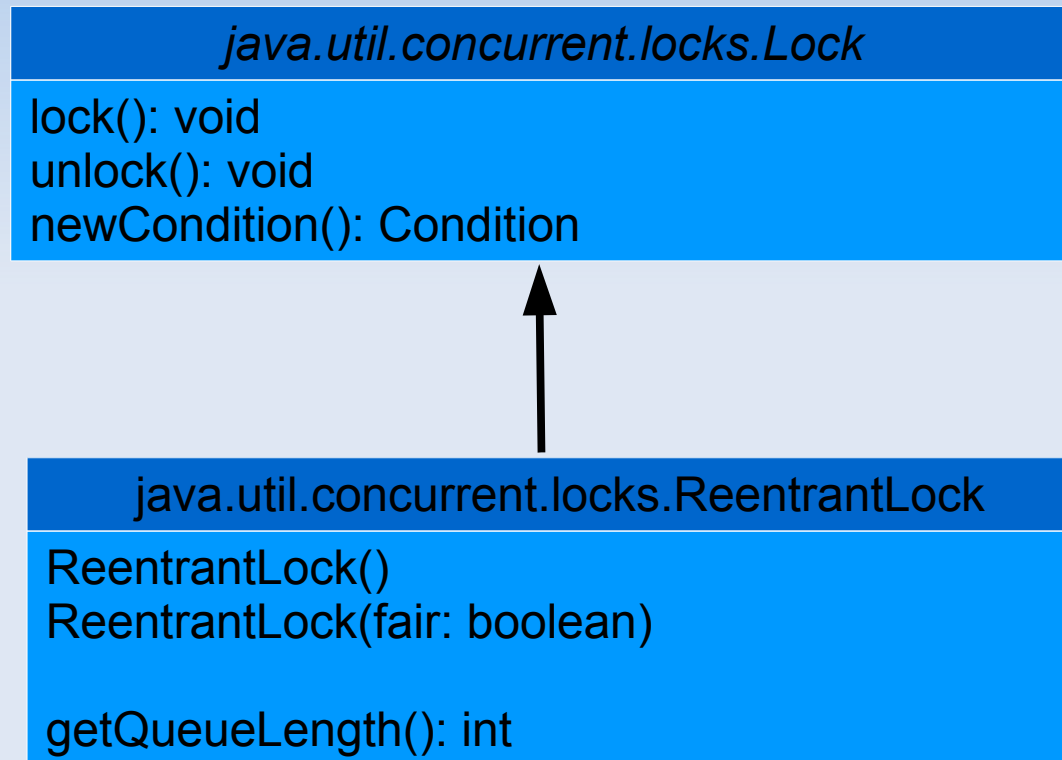
# Thread Synchronization

- This results in the exact same outcome as synchronizing each method.
- The thread attempts to lock the account object, then runs a method on it and releases it.
- If we have many different objects and only one needed to be locked, then we can use this synchronized block to just lock the one object we need and complete our task.
- This would improve the overall runtime of our multi-threaded programs.

# Synchronization using locks

- Using a synchronized block, or declaring our method as synchronized will implicitly place a lock on the object being used.
- However, Java allows you to explicitly create locks to use in your code.
- All locks in Java implement the `java.util.concurrent.locks.Lock` interface.
- We will be looking at the `ReentrantLock` for our locking needs.

# Synchronization using locks



The `ReentrantLock(boolean)` constructor allows you to specify a fairness policy. The fairness policy states that the longest waiting thread will get the lock next.

# Synchronization using locks

```
import java.util.concurrent.locks.*;

public class Account{
    private Lock lock = new ReentrantLock();
    private int balance = 0;

    public int getBalance(){
        return balance;
    }

    public void deposit(int amount){
        lock.lock();
        try{
            balance = balance + amount;
        }finally{
            lock.unlock();
        }
    }
}
```

Its good practice to always immediately follow a lock statement with a try-catch-finally, and put the unlock call in the finally block. This guarantees it will always run and not cause a deadlock. Even if the statement doesn't throw an exception, you should use a try finally block.

# Synchronization using locks

```
public class UseAccount{
    private static Account account = new Account();

    public static void main(String[] args){
        Thread husband = new Thread(new Person());
        Thread wife = new Thread(new Person());

        husband.start();
        wife.start();
        try{
            husband.join();
            wife.join();
        }catch(InterruptedException ignore){}

        System.out.println("The account balance is: " +
                           account.getBalance());
    }

    static class Person implements Runnable{
        public void run(){
            for (int i = 0; i < 10000; i++){
                account.deposit(1);
            }
        }
    }
}
```



# Synchronization using locks

- Although using a lock in this instance is overkill, there are other things that locks can provide that a synchronized block simply can't.

```
java.util.concurrent.locks.ReentrantLock  
getHoldCount(): int  
getOwner: Thread  
getQueuedThreads(): Collection  
getQueueLength(): int  
getWaitingThreads(): Collection  
hasQueuedThreads(): boolean  
hasQueuedThread(t: Thread): boolean  
tryLock(): boolean  
tryLock(timeout: long, unit: TimeUnit): boolean
```

# Communication between threads

- Locks can also be used to provide a way for threads to cooperate.
- The `Lock.newCondition()` method will create a cooperation condition bound to the lock.

*java.util.concurrent.Condition*

`await(): void`  
`signal(): void`  
`signalAll(): void`

# Communication between threads

```
import java.util.concurrent.locks.*;

public class Account{
    private Lock lock = new ReentrantLock();
    private Condition depositCondition = lock.newCondition();
    private int balance = 0;

    public int getBalance(){
        return balance;
    }

    public void deposit(int amount){
        lock.lock();
        try{
            System.out.println("Deposit made: " + amount);
            balance += amount;
            depositCondition.signalAll();
        }finally{
            lock.unlock();
        }
    }
}
```

# Communication between threads

```
public void withdraw(int amount){
    lock.lock();
    try{
        while(balance < amount){
            System.out.println("\t\t\tWait for deposit");
            depositCondition.await();
        }
        balance -= amount;
        System.out.println("\t\t\tWithdraw made: " + amount);
    }catch(InterruptedException ignore){

    }finally{
        lock.unlock();
    }
}
```

# Communication between threads

```
public class UseAccount{
    private static Account account = new Account();

    public static void main(String[] args){
        Thread deposit = new Thread(new DepositTask());
        Thread withdraw = new Thread(new WithdrawTask());

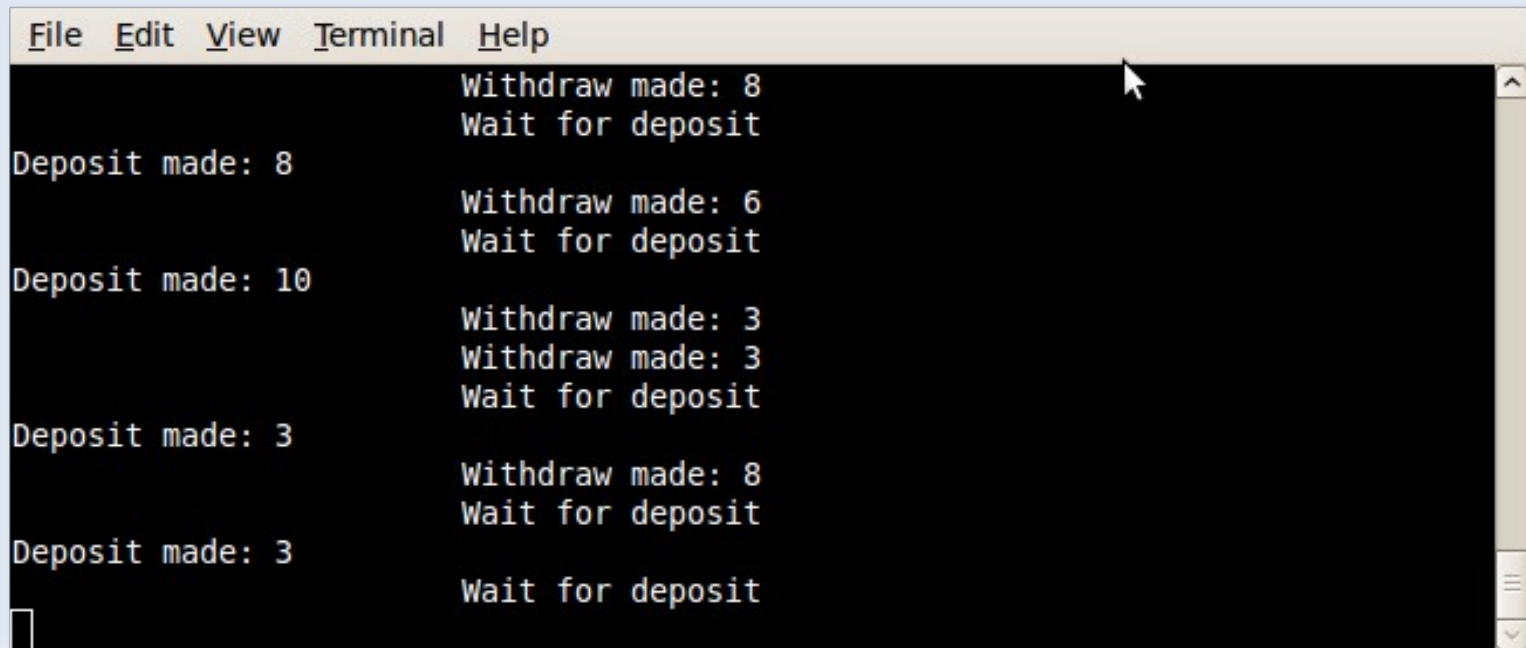
        deposit.start();
        withdraw.start();
        try{
            deposit.join();
            withdraw.join();
        }catch(InterruptedException ignore){}

        System.out.println("The account balance is: " + account.getBalance());
    }

    static class DepositTask implements Runnable{
        public void run(){
            try{
                for (int i = 0; i < 10000; i++){
                    account.deposit((int)(Math.random() * 10) + 1);
                    Thread.sleep(300); //let the withdraw have a shot
                }
            }catch(InterruptedException ignore){}
        }
    }
}
```

# Communication between threads

```
static class WithdrawTask implements Runnable{  
    public void run(){  
        try{  
            for (int i = 0; i < 10000; i++){  
                account.withdraw((int) (Math.random() * 10) + 1);  
                Thread.sleep(100);  
            }  
        }catch (InterruptedException ignore){}  
    }  
}
```



The screenshot shows a terminal window with a menu bar (File, Edit, View, Terminal, Help) and a black background with white text. The output is interleaved, showing the execution of two threads. The 'Deposit' thread outputs 'Deposit made: 8', 'Deposit made: 10', 'Deposit made: 3', and 'Deposit made: 3'. The 'Withdraw' thread outputs 'Withdraw made: 8', 'Withdraw made: 6', 'Withdraw made: 3', 'Withdraw made: 3', and 'Withdraw made: 8'. Each withdrawal is followed by 'Wait for deposit', indicating a synchronization point where the thread waits for a deposit to be made before proceeding.

```
File Edit View Terminal Help  
Withdraw made: 8  
Wait for deposit  
Deposit made: 8  
Withdraw made: 6  
Wait for deposit  
Deposit made: 10  
Withdraw made: 3  
Withdraw made: 3  
Wait for deposit  
Deposit made: 3  
Withdraw made: 8  
Wait for deposit  
Deposit made: 3  
Wait for deposit
```

# Using threads instead of timers

- You can use a separate thread instead of a Timer object to control animation.
- Threads are better at giving you more precise timing.
- Timers use less resources but can provide bad performance if the actionPerformed method takes a long time to run.

# Using threads instead of timers

```
import javax.swing.*;

public class FlashText extends JFrame{
    private JLabel label = new JLabel("Flashing Text");
    private FlashTimer timer = new FlashTimer(50);

    public FlashText(){
        label.setHorizontalAlignment(JLabel.CENTER);
        add(label);
        new Thread(timer).start();
    }

    public static void main(String[] args){
        JFrame frame = new FlashText();
        frame.setSize(300,300);
        frame.setLocationRelativeTo(null);
        frame.setDefaultCloseOperation(JFrame.EXIT_ON_CLOSE);
        frame.setVisible(true);
    }
}
```



# Using threads instead of timers

```
private class FlashTimer implements Runnable{
    private int waitTime = 100;
    public FlashTimer(){

    }

    public FlashTimer(int waitTime){
        this.waitTime = waitTime;
    }

    public void run(){
        boolean on = true;
        try{
            while(true){
                if (!on){
                    label.setText("Flashing Text");
                    on = true;
                }else{
                    label.setText(null);
                    on = false;
                }

                Thread.sleep(waitTime);
            }
        }catch (InterruptedException ignore){}
    }
}
```

# Lab Assignment

No lab this week (The exercises in this chapter suck!)

Homework

- Page 972 # 29.5 (Running fans)

# Acknowledgments

Introduction to Java Programming by Y. Daniel Liang

